

Optimal Portfolio Weights IV: Constrained Optimization

```
In [1]: # Working with data:
import numpy as np                                # For scientific computing
import pandas as pd                               # Working with tables.

# Downloading files:
import requests, zipfile, io                      # To access websites

# Specific data providers:
from tiingo import TiingoClient                   # Stock prices.
import quandl                                     # Economic data, futures p

# API keys:
tiingo = TiingoClient({'api_key': 'xxxx'})
quandl.ApiConfig.api_key = 'yyyy'

# Plotting:
import matplotlib.pyplot as plt                   # Basic plot library.
plt.style.use('ggplot')                           # Make plots look nice
```

SPRD sector ETFs

```
In [2]: sectors = ['XLY', 'XLP', 'XLE', 'XLF', 'XLV', 'XLI', 'XLB', 'XLK', 'XLU']

PRICE = tiingo.get_dataframe(sectors, '1999-1-1', metric_name='adjClose')
PRICE.index = pd.to_datetime(PRICE.index).tz_convert(None)
RET = PRICE.pct_change()

RATES = quandl.get(['FRED/FEDFUNDS', 'FRED/DGS1']) / 100
RATES.columns = ['FedFunds', 'Treasury_1']
```

```
In [3]: PRICE.plot(logy=True)
```

Out[3]: <AxesSubplot:>



Get SPY as benchmark:

```
In [4]: spy = tiingo.get_dataframe(['SPY'], '1999-1-1', metric_name='adjClose')
```

```
spy.index = pd.to_datetime(spy.index).tz_convert(None)
spy[-3:]
```

Out[4]:

| | SPY |
|------------|--------|
| 2021-04-12 | 411.64 |
| 2021-04-13 | 412.86 |
| 2021-04-14 | 411.45 |

Backtest function:

In [5]:

```
def get_rebalance_dates(frequency, start_date):
    price = PRICE[PRICE.index>start_date]
    group = getattr(price.index, frequency)
    return price[:1].index.union(price.groupby([price.index.year, group]).tail(1)

def compare_performance(t):
    t.add(1).cumprod().plot()
    t.add(1).cumprod().plot(logy=True)

    annual_returns = t[:'2020'].add(1).resample('A').prod().sub(1)
    r_annual_Tbill = RATES.Treasury_1.resample('A').first()

    x = pd.DataFrame()
    x['Average_returns'] = annual_returns.mean()
    x['Geometric_average'] = annual_returns.add(1).prod().pow(1/len(annual_returns)).sub(1)
    x['Risk_premium'] = annual_returns.sub(r_annual_Tbill, 'rows').dropna()
    x['Volatility'] = t[:'2020'].std() * 252**0.5
    x['Sharpe_ratio'] = x.Risk_premium / x.Volatility

    return x

def run_backtest(frequency, backtest_start='1900-1-1'):

    rebalance_dates = get_rebalance_dates(frequency, backtest_start)

    portfolio_value = pd.Series(1, index=rebalance_dates)
    weights = pd.DataFrame(columns=PRICE.columns, index=rebalance_dates)
    trades = pd.DataFrame(columns=PRICE.columns, index=rebalance_dates)

    previous_positions = weights.iloc[0]

    for i in range(1, len(rebalance_dates)-1):
        start_date = rebalance_dates[i]
        end_date = rebalance_dates[i+1]

        cum_ret = RET[start_date:end_date][1:].add(1).cumprod()

        start_weights = select_weights(start_date) # Call "select_weights()"

        new_positions = portfolio_value.iloc[-1] * start_weights

        start_to_end_positions = new_positions * cum_ret
```

```

start_to_end_value      = start_to_end_positions.sum('columns')

portfolio_value = portfolio_value.append(start_to_end_value)

weights = weights.append(start_to_end_positions.div(start_to_end_value, 'columns'))

trades.loc[start_date] = new_positions - previous_positions
previous_positions      = start_to_end_positions.iloc[-1]      # Previous

return portfolio_value.pct_change(), weights, trades

```

Backtest strategy that equal-weights all assets:

```

In [6]: pd.Series(1/len(sectors), index=sectors) # put 1/n in each asset

```

```

Out[6]: XLY      0.111111
        XLP      0.111111
        XLE      0.111111
        XLF      0.111111
        XLV      0.111111
        XLI      0.111111
        XLB      0.111111
        XLK      0.111111
        XLU      0.111111
        dtype: float64

```

```

In [7]: def select_weights(date):
        return pd.Series(1/len(sectors), index=sectors) # Equal weights (1/n)

ew, weights, trades = run_backtest('month', '2000-1-1')
ew = ew.rename('EW')

t = pd.DataFrame(ew).join(spy.pct_change())
compare_performance(t)

```

```

Out[7]:

```

| | Average_returns | Geometric_average | Risk_premium | Volatility | Sharpe_ratio |
|------------|-----------------|-------------------|--------------|------------|--------------|
| EW | 0.091669 | 0.078461 | 0.072226 | 0.189818 | 0.380502 |
| SPY | 0.082703 | 0.067384 | 0.063260 | 0.198243 | 0.319105 |





Minimum volatility weights:

```
In [9]: def select_weights(date):
cov      = RET[:date][-100:].cov() * 252 # Use most recent 100 returns up to
cov_inv = pd.DataFrame(np.linalg.inv(cov), columns=cov.columns, index=cov.in
w = cov_inv.sum() / cov_inv.sum().sum() # Minimum-volatility portfo
return w

mv, weights, trades = run_backtest('month', '2000-1-1')
mv = mv.rename('MinVol')

t = pd.DataFrame(mv).join(ew).join(spy.pct_change())
compare_performance(t)
```

| | Average_returns | Geometric_average | Risk_premium | Volatility | Sharpe_ratio |
|---------------|-----------------|-------------------|--------------|------------|--------------|
| MinVol | 0.082771 | 0.075506 | 0.063328 | 0.144292 | 0.438889 |
| EW | 0.091669 | 0.078461 | 0.072226 | 0.189818 | 0.380502 |
| SPY | 0.082703 | 0.067384 | 0.063260 | 0.198243 | 0.319105 |






Plot the weights:

```
In [10]: weights.plot()
```

Out[10]: <AxesSubplot:>



Note how these weights contain short positions (the negative weights).

Why short positions? All these ETFs have relatively high positive correlation, and so the min-vol portfolio requires that we short some assets:

```
In [11]: RET.corr()
```

```
Out[11]:
```

| | XLY | XLP | XLE | XLF | XLV | XLI | XLB | XLK | XL |
|------------|----------|----------|----------|----------|----------|----------|----------|----------|---------|
| XLY | 1.000000 | 0.630802 | 0.550253 | 0.757101 | 0.698223 | 0.818486 | 0.719942 | 0.721419 | 0.51033 |
| XLP | 0.630802 | 1.000000 | 0.477009 | 0.578416 | 0.619822 | 0.634423 | 0.564218 | 0.490179 | 0.60084 |
| XLE | 0.550253 | 0.477009 | 1.000000 | 0.579339 | 0.510408 | 0.661278 | 0.688992 | 0.474144 | 0.51416 |
| XLF | 0.757101 | 0.578416 | 0.579339 | 1.000000 | 0.632924 | 0.786996 | 0.690217 | 0.630460 | 0.50337 |
| XLV | 0.698223 | 0.619822 | 0.510408 | 0.632924 | 1.000000 | 0.714460 | 0.604808 | 0.678747 | 0.54092 |
| XLI | 0.818486 | 0.634423 | 0.661278 | 0.786996 | 0.714460 | 1.000000 | 0.809604 | 0.729945 | 0.55670 |
| XLB | 0.719942 | 0.564218 | 0.688992 | 0.690217 | 0.604808 | 0.809604 | 1.000000 | 0.588904 | 0.50931 |
| XLK | 0.721419 | 0.490179 | 0.474144 | 0.630460 | 0.678747 | 0.729945 | 0.588904 | 1.000000 | 0.46187 |

| | XLY | XLP | XLE | XLF | XLV | XLI | XLB | XLK | XL |
|------------|----------|----------|----------|----------|----------|----------|----------|----------|---------|
| XLU | 0.510337 | 0.600842 | 0.514161 | 0.503379 | 0.540921 | 0.556703 | 0.509319 | 0.461872 | 1.00000 |

Problem: shorting is costly:

- we have to pay a borrow fee when we borrow stocks
- we have to keep the (negative) value of the short position in cash (as collateral)
- we earn interest on the collateral (benchmark rate minus some spread; currently this interest is zero)
- we cannot use the collateral to finance long positions

Example:

initially you have \$100 cash in your account (and nothing else),
now you short an asset and receive 50 in cash from the sale and then you buy another asset for 150.

intial account:

- cash: 100
- equity: 100 (account value)

After the transction:

- stock A: -50 (short, pay fee)
- stock B: 150 (long)
- collateral: 50 (from short sale, earns interest)
- margin loan: 50 (to pay for stock B, pay interest)
- cash total: 0
- equity: 100 (account value)

(The interest you pay on the margin loan is higher than the interest you earn on the collateral (currently you earn zero interest, and pay about 1.5%))

Most recent weights of the min-vol portfolio:

```
In [12]: weights.iloc[-1] # Last row of weights table
```

```
Out[12]: XLY      0.072929
XLP      0.530791
XLE     -0.042154
XLF      0.066804
XLV      0.486284
XLI      0.200180
XLB     -0.223202
XLK     -0.053567
XLU     -0.038065
Name: 2021-04-14 00:00:00, dtype: float64
```

"Clip" weights at zero:

```
In [13]: weights.iloc[-1].clip(0)
```

```
Out[13]: XLY      0.072929
          XLP      0.530791
          XLE      0.000000
          XLF      0.066804
          XLV      0.486284
          XLI      0.200180
          XLB      0.000000
          XLK      0.000000
          XLU      0.000000
          Name: 2021-04-14 00:00:00, dtype: float64
```

Sum of these weights:

```
In [14]: weights.iloc[-1].clip(0).sum()
```

```
Out[14]: 1.3569883034998693
```

→ when we set negative weights to zero, the weights don't sum to one anymore!

How can we transform the weights so that they sum to one?

→ divide the clipped weights by their sum:

```
In [15]: w = weights.iloc[-1].clip(0) / weights.iloc[-1].clip(0).sum()
          w
```

```
Out[15]: XLY      0.053744
          XLP      0.391153
          XLE      0.000000
          XLF      0.049229
          XLV      0.358355
          XLI      0.147518
          XLB      0.000000
          XLK      0.000000
          XLU      0.000000
          Name: 2021-04-14 00:00:00, dtype: float64
```

Check the sum:

```
In [16]: w.sum()
```

```
Out[16]: 1.0
```

Add this to our select_weights function:



```
In [18]: def select_weights(date):
          cov      = RET[:date][-100:].cov() * 252 # Use most recent 100 returns up to
          cov_inv  = pd.DataFrame(np.linalg.inv(cov), columns=cov.columns, index=cov.in
          w = cov_inv.sum() / cov_inv.sum().sum()    # Minimum-volatility portfo
          w = w.clip(0) / w.clip(0).sum()           # Clip weights at zero
          return w

          mv_clipped, weights, trades = run_backtest('quarter', '2000-1-1')
          mv_clipped = mv_clipped.rename('All_MinVol_Clipped')

          t = pd.DataFrame(mv).join(mv_clipped).join(ew).join(spy.pct_change())
          compare_performance(t)
```

Out[18]:


| | Average_returns | Geometric_average | Risk_premium | Volatility | Sharpe_ratio |
|---------------------------|-----------------|-------------------|--------------|------------|--------------|
| MinVol | 0.082771 | 0.075506 | 0.063328 | 0.144292 | 0.438889 |
| All_MinVol_Clipped | 0.087075 | 0.078038 | 0.067633 | 0.155712 | 0.434345 |
| EW | 0.091669 | 0.078461 | 0.072226 | 0.189818 | 0.380502 |
| SPY | 0.082703 | 0.067384 | 0.063260 | 0.198243 | 0.319105 |



Note how the unconstrained min-vol portfolio has a lower volatility than the "clipped" portfolio (in the table above).

```
In [19]: weights.plot()
```

```
Out[19]: <AxesSubplot:>
```

Here we chose weights like this:

1. find minimum volatility weights
2. set all negative weights to zero
3. recalibrate weights so they sum to zero: divide all weights by their sum

But this procedure is only an approximation of the optimal weights.

To find the actual weights we need to:

- minimize volatility under the constraint that weights are ≥ 0

Install `cvxopt` library for constrained optimization:

In [20]:

```
pip install CVXOPT
```

Requirement already satisfied: CVXOPT in /Users/janschneider/opt/anaconda3/lib/python3.7/site-packages (1.2.6)

Note: you may need to restart the kernel to use updated packages.

In [21]:

```
import cvxopt
from cvxopt import matrix, solvers
solvers.options['show_progress'] = False
```

Now use the `cvxopt` like this:

In [22]:

```
minimum_weight = 0    # or try for example 0.05 to get minimum diversification

cov = RET.cov()

n = len(cov)
P = matrix( cov.values )
q = matrix( np.zeros(n) )
G = matrix( -np.identity(n) )
h = matrix( -np.ones(n)*minimum_weight)
A = matrix( np.ones(n), (1,n) )           # weights sum to 1
b = matrix(1.0)

sol = solvers.qp(P, q, G, h, A, b)

pd.Series({cov.index[i]:sol['x'][i] for i in range(n)})
```

```
Out[22]: XLY      0.000190
          XLP      0.608019
          XLE      0.000146
          XLF      0.000010
          XLV      0.224365
          XLI      0.000241
          XLB      0.000618
          XLK      0.000346
          XLU      0.166064
          dtype: float64
```

Compare clipped weights to constrained minimum:

```
In [23]: weights = pd.DataFrame()

# Min-vol with short
cov      = RET.cov()
cov_inv  = pd.DataFrame(np.linalg.inv(cov), columns=cov.columns, index=cov.index)
w        = cov_inv.sum() / cov_inv.sum().sum()

weights['With_shortselling'] = w

# Min-vol, clipped at zero:
cov      = RET.cov()
cov_inv  = pd.DataFrame(np.linalg.inv(cov), columns=cov.columns, index=cov.index)
w        = cov_inv.sum() / cov_inv.sum().sum()
w        = w.clip(0) / w.clip(0).sum()

weights['Clipped'] = w

# Constrained minimum:
minimum_weight = 0
cov = RET.cov()

n = len(cov)
P = matrix( cov.values )
q = matrix( np.zeros(n) )
G = matrix( -np.identity(n) )
h = matrix( -np.ones(n)*minimum_weight)
A = matrix( np.ones(n), (1,n))
b = matrix(1.0)

sol = solvers.qp(P, q, G, h, A, b)

weights['Constrained'] = pd.Series({cov.index[i]:sol['x'][i] for i in range(n)})
weights
```

```
Out[23]:
```

| | With_shortselling | Clipped | Constrained |
|-----|-------------------|----------|-------------|
| XLY | 0.022612 | 0.018936 | 0.000190 |
| XLP | 0.622693 | 0.521466 | 0.608019 |
| XLE | -0.016874 | 0.000000 | 0.000146 |
| XLF | -0.176690 | 0.000000 | 0.000010 |

| | With_shortselling | Clipped | Constrained |
|------------|-------------------|----------|-------------|
| XLV | 0.276966 | 0.231941 | 0.224365 |
| XLI | 0.048347 | 0.040488 | 0.000241 |
| XLB | 0.040495 | 0.033912 | 0.000618 |
| XLK | -0.000558 | 0.000000 | 0.000346 |
| XLU | 0.183008 | 0.153258 | 0.166064 |


Put the constrained optimization code into the select weights function:


```
In [24]: def select_weights(date):
cov = RET[:date][-100:].cov()
minimum_weight = 0
n = len(cov)
P = matrix( cov.values )
q = matrix( np.zeros(n) )
G = matrix( -np.identity(n) )
h = matrix( -np.ones(n)*minimum_weight )
A = matrix( np.ones(n), (1,n) )
b = matrix(1.0)
sol = solvers.qp(P, q, G, h, A, b)
return pd.Series({cov.index[i]:sol['x'][i] for i in range(n)})

mv_cvxopt, weights, trades = run_backtest('month', '2000-1-1')
mv_cvxopt = mv_cvxopt.rename('MinVol_cvxopt')

t = pd.DataFrame(mv).join(mv_clipped).join(mv_cvxopt).join(ew).join(spy.pct_chan
compare_performance(t)
```

| | Average_returns | Geometric_average | Risk_premium | Volatility | Sharpe_ratio |
|---------------------------|-----------------|-------------------|--------------|------------|--------------|
| MinVol | 0.082771 | 0.075506 | 0.063328 | 0.144292 | 0.438889 |
| All_MinVol_Clipped | 0.087075 | 0.078038 | 0.067633 | 0.155712 | 0.434345 |
| MinVol_cvxopt | 0.074977 | 0.066212 | 0.055534 | 0.146711 | 0.378528 |
| EW | 0.091669 | 0.078461 | 0.072226 | 0.189818 | 0.380502 |
| SPY | 0.082703 | 0.067384 | 0.063260 | 0.198243 | 0.319105 |





Note how the coonstrained minimum portfolio has a lower volatility than the unconstrained min-vol portfolio that was clipped at zero (in the table above).